Traffic Flow Prediction

Using Graph Neural Networks

Jahnavi Pokala

jpokala@kent.edu

811316337

Abstract

Traffic flow prediction is an essential component of intelligent transportation systems, enabling traffic management, congestion reduction, and optimized urban mobility. Traditional methods, such as time-series models and early deep learning techniques, have had limited success but typically cannot cope with the complex spatial-temporal correlations of traffic networks satisfactorily. Modern Graph Neural Network (GNN) developments offer a solution with the potential to depict traffic networks in graph form, with intersections or road segments treated as nodes and their interactions defined as edges. The present development of traffic flow forecasting using GNNs has been reviewed here at length to cover the recently emerged top-performing models DCRNN, STGCN, Graph WaveNet, ASTGCN, and STDGAT. We explain how they enhance spatial-temporal learning and evaluate their performance on real-world data. We also introduce industry applications in navigation services, smart city infrastructure, and autonomous driving. Data sparsity, scalability, model interpretability, and privacy issues are addressed, as well as potential future research directions. This study concludes that GNNs are a revolutionary method of enhancing traffic forecasting technology in ever more complex urban environments.

1. Introduction

Traffic congestion is, maybe, the largest problem of modern urban culture. As there has been a growing population in the cities, an immense demand has been generated to have efficient traffic management systems. Traffic flow prediction, where traffic conditions in the future are forecasted using past as well as current data, is the fundamental functionality of intelligent transport systems (ITS). Accurate traffic forecasting could, in theory, lead to better traffic management, reduced congestion, reduced carbon emissions, increased public safety, and more effective urban planning. Traditional traffic forecasting methods, while beneficial to some degree, are not, by and large, in a position to reflect the complex, dynamic, and interconnected nature of real traffic systems.

The last few years have seen monumental advances in machine learning, i.e., deep learning, which have completely revolutionized traffic prediction. These early models, such as time-series-based statistical models (e.g., ARIMA) and recurrent neural networks (e.g., LSTM), were significantly superior to the conventional techniques. These don't, however, function very effectively to leverage the spatial interdependencies of traffic networks where conditions at a point heavily depend on nearby points. In place of this absence, researchers utilized Graph Neural Networks (GNNs), a deep learning network designed specifically to handle graph-structured data.

Traffic networks naturally lend themselves to graph representation, where sensors, intersections, or road segments are represented as nodes and interconnectivity as edges. GNNs leverage such structure to represent complex spatial relationships as neighboring node information aggregation. When combined with temporal evolutions of spatial attributes, GNN-based models have shown better traffic flow prediction compared to traditional and other deep learning methods.

The goal of this paper is to explore recent developments in traffic flow prediction with Graph Neural Networks. We seek to present an extensive overview of research, point out effective models and methods, outline practical applications of the provided models in real-world scenarios, and mark the limitations and challenges of the current systems. In addition, we will discuss potential future directions that can further increase the capability of GNN-based traffic forecasting models such that they can respond to the changing needs of smart cities. Through this analysis, our objective is to illustrate the pivotal role that GNNs can assume in developing traffic management technologies and urban mobility futures.

2. Background and Motivation

Accurate prediction of traffic flow has always been a dream in transportation engineering and urban planning. The early methods, which were largely mathematical modeling and statistical, formed the basis for traffic explanation and prediction. Some of the initial methods, such as the Auto-Regressive Integrated Moving Average (ARIMA) model and its extensions, treated traffic data as a time series using past data for predicting future traffic conditions. These models were easy to implement and comprehend, but suffered the severe disadvantage of being incapable of accommodating non-linear relationships as well as spatial relationships among components of a traffic network.

As machine learning unfolded its wings extensively, newer and more powerful models like Support Vector Machines (SVM) and shallow neural networks became the talk of the town, which were more accurate at learning higher-order structures in data. However, the true break came when deep learning evolved with the help of Recurrent Neural Networks (RNNs) and their gateway relatives like Long Short-Term Memory (LSTM) networks. Such methods prevailed in acquiring temporal dependences, quite decently simulating sequences of traffic patterns over time. Convolutional Neural Networks (CNNs), initially devised for image processing, were utilized in processing space patterns of traffic data by handling traffic maps as grid-based pictures.

Despite all these advances, there was still one essential limitation. Traffic networks are inherently non-Euclidean — they're best described not as grids, but as serrated networks where roads (edges) and intersections (nodes) create complicated, dynamic topologies. Grid-based models such as CNNs can't natively encode interconnectivity between far-flung yet linked nodes in a traffic network, thus losing precious spatial information. Alternatively, although LSTM models are effective with time sequences, they have naíve spatial data handling and generally take spatial independence or naive distance-based assumptions.

This restriction resulted in the investigation of Graph Neural Networks (GNNs) for making traffic flow predictions. GNNs are particularly well-equipped to handle graph-structured data, such that models can learn from node features themselves as well as from graph structure. In traffic networks, GNNs can learn the intricate spatial correlations where traffic at one node influences and is influenced by its neighboring nodes. Additionally, by the incorporation of temporal modeling methods, scientists were even able to create spatio-temporal GNN models, which are capable of learning how traffic changes over space and time.

The interest in working on GNNs for traffic forecasting is thus twofold: first, to overcome the spatial pattern limitation of the conventional as well as deep learning methods; and second, to enable more precise, scalable, and robust forecasting of traffic in more and more complex urban scenarios. As urban cities are growing towards smart infrastructure, there will be more and more demand for high-precision real-time traffic forecasting systems. GNNs, in that they can handle large-scale, dynamic, and irregular data, offer a solution worthy of meeting these shifting needs. In addition to having the potential to make traffic more efficient, their use has the potential to add significantly to city sustainability, economic productivity, and public safety.

3. Understanding Graph Neural Networks

The capacity for copying and learning based on ordered relational data has increasingly become inescapable in most applications, particularly in traffic systems where the relationship between different points dynamically changes and comprises maximum complexity. Conventional deep learning architectures such as Recurrent Neural Networks (RNNs) and Convolutional Neural Networks (CNNs) are highly effective with grid-structured data (e.g., images and sequences) but not easily transferable in the case of irregular and non-Euclidean data structures such as road networks. Graph Neural Networks (GNNs) have bridged the gap through direct graph-structured data processing (Bronstein et al., 2017; Wu et al., 2020).

A graph G (V, E) is a set of vertices (or nodes) V and edges E among the vertices. In traffic networks, nodes will likely represent elements such as road junctions, sensors, or areas, and edges represent roads or connectivity between sites. Nodes and edges can each be labeled by feature vectors, such as traffic volume, speed, and road type.

GNNs are calculated using a procedure known as message passing, wherein each node learns from its neighbor nodes and calculates its state on the basis of these. By iteratively performing this procedure, the model can learn complex spatial relations of the graph. A node gathers neighbor feature information in each layer of a GNN and performs a transformation (in most cases, a neural network) to calculate the new feature representation (Kipf & Welling, 2017).

There have been different GNN architectures suggested, differing in update and aggregation operations:

- Graph Convolutional Networks (GCNs): They utilize the graph convolutional concept with a weighted average of the neighboring nodes' features and a non-linear function (Kipf & Welling, 2017).
- Graph Attention Networks (GATs): Instead of using the same weights across the board, GATs learn to use different importances to different neighbors with an attention mechanism (Veliković et al., 2018).
- GraphSAGE: Rather than calculating representations for an entire batch of nodes simultaneously, GraphSAGE samples and pools features from a fixed-sized neighborhood of neighbors so that it can be applied to large graphs (Hamilton et al., 2017).

Spatial interaction among various road segments is relevant in traffic flow forecasting. An example is when a congestion event at one intersection has direct effects on surrounding intersections. Such spatial interaction is inherent in GNNs, where information can propagate from nearby nodes in the graph of the traffic network. Moreover, traffic data are time-varying — traffic behavior changes as time goes on due to events such as rush hour, weather, and accidents. Hence, more recent models cause GNNs to incorporate temporal modeling methods such as RNNs, LSTMs, or temporal convolutions to create Spatio-Temporal Graph Neural Networks (Yu et al., 2018; Li et al., 2018).

Some of the highest performers include the Diffusion Convolutional Recurrent Neural Network (DCRNN) (Li et al., 2018), where traffic flow is defined as a diffusion on a directed graph. The other is the Spatio-Temporal Graph Convolutional Network (STGCN), where graph convolutions are utilized alongside temporal convolutional layers in trying to handle spatial and temporal dependencies simultaneously (Yu et al., 2018).

The model's decision to graph itself is significant, too. There are some models that are constructed on a static graph of physical roads, while there are others that dynamically change the graph topology adaptively as per changing traffic patterns (Cui et al., 2019). Graph construction in this way leads the model to learn as per changing traffic patterns, such as road closures or ad-hoc congestion events.

In general, GNNs give a highly flexible and robust traffic flow forecasting model by learning spatial patterns and dynamic temporal patterns characteristic of traffic systems very well. Their capability to generalize over large-scale and irregular networks makes them best suited for application in real-world intelligent transportation systems. With more studies, newer and more advanced variants of GNNs are being worked upon with features like attention mechanisms, dynamic graphs, and multi-scale representations to push the limits of what traffic prediction models are capable of doing.

4. Literature Review

Traffic flow forecasting with Graph Neural Networks (GNNs) has been one of the prominent topics of vast amounts of research work in recent years, giving rise to an abundance of new models. We present the best state-of-the-art methods that have contributed to getting us where we are today, describing their approach, contribution, and shortcomings.

4.1 Diffusion Convolutional Recurrent Neural Network (DCRNN)

Diffusion Convolutional Recurrent Neural Network (DCRNN) is a seminal work of Li et al. (2018) and a new model integrating graph space modeling with temporal dynamics to predict traffic. DCRNN represents traffic flow as directed graph diffusion in which information diffuses in different directions throughout the network. Diffusion convolution operations in a sequence-to-sequence network of Gated Recurrent Units (GRUs) form this model.

Li et al. showed that DCRNN successfully outperformed baselines like Historical Average, ARIMA, and vanilla LSTM on real traffic data like METR-LA and PEMS-BAY. Through explicit spatial diffusion modeling and temporal correlation, DCRNN realized more than 12% improvement in mean absolute error (MAE) compared to non-graph algorithms. DCRNN, however, relies on a fixed graph pattern, which could restrict its adaptability to learn from changing traffic patterns.

4.2 Spatio-Temporal Graph Convolutional Network (STGCN)

Yu et al. (2018) proposed the Spatio-Temporal Graph Convolutional Network (STGCN), which integrates graph convolutions and temporal convolutions into a purely convolutional architecture that removes the use of recurrent networks. STGCN interweaves graph convolutional layers (for spatial information) and temporal convolutional layers (for time-series information).

On METR-LA data, STGCN outperformed the conventional RNN-based approaches and was much more efficient to train since it did not include the recurrent units. Although STGCN uses a static adjacency matrix and thus is less sensitive to abrupt breaks like road closures or accidents, its efficiency was significantly enhanced by not having to struggle with the recurrent units.

4.3 Graph WaveNet

Graph WaveNet (Wu et al., 2019) also promoted a more adaptive approach by training the adjacency matrix instead of predefining it. It combines dilated temporal convolutions and graph convolutions and incorporates adaptive adjacency matrices able to learn intrinsic spatial relations.

METR-LA and PEMS-BAY experiments revealed that Graph WaveNet outperformed STGCN and DCRNN, particularly for long-term traffic forecasting tasks. Its capacity to learn dynamic graphs allows it to capture evolving traffic patterns. But with the model's complexity comes a higher requirement for computation, hence less desirable when applied to low-end hardware.

4.4 Attention-Based Spatial-Temporal Graph Convolutional Network (ASTGCN)

Guo et al. (2019) introduced the Attention-Based Spatial-Temporal Graph Convolutional Network (ASTGCN) to further improve model interpretability and performance. ASTGCN integrates both the spatial and temporal attention mechanisms to enable the model to concentrate on the most significant nodes and steps in time for value prediction.

This attention mechanism improves the quality of predictions as well as model interpretability, allowing better interpretation of the traffic dynamics' influencing factors. ASTGCN has achieved state-of-the-art results on benchmark datasets but needs to be enforced through extensive hyperparameter fine-tuning of several attention parameters, which hinders model training.

4.5 Spatio-Temporal Dynamic Graph Attention Network (STDGAT)

Latest progress saw the publication of the Spatio-Temporal Dynamic Graph Attention Network (STDGAT) (Song et al., 2020), which dynamically builds graphs based on real-time traffic conditions and applies attention in both spatial and temporal directions.

STDGAT performs very efficiently, especially with time-varying traffic such as weather alteration or special activities. STDGAT dynamically builds graph connectivity at inference and hence is very resilient, but with more computational complexity.

4.6 Comparative Analysis

Model	Graph Type	Attention	Strengths	Limitations
DCRNN	Static	No	Strong spatio-temporal modeling	Limited flexibility
STGCN	Static	No	Fast training, simple	Fixed adjacency matrix
Graph WaveNet	Dynamic (learned)	No	Captures hidden dependencies	High complexity
ASTGCN	Static	Yes (both spatial and temporal)	High interpretability	Heavy hyperparameter tuning
STDGAT	Dynamic	Yes	Adapts to real-time traffic	Computationally expensive

4.7 Limitations Observed in Existing Research

Although there is tremendous progress, there are some challenges still awaiting solution:

- Dynamic graphs do not automatically adjust to varying traffic unless dynamically learned.
- Excessive computational expense of sophisticated models such as Graph WaveNet and STDGAT.
- Interpretability: Most models are black-box in nature, and hence hard to trust for safety-critical use.
- Availability of data: High definition and real-time traffic data are not always available, especially in developing countries.

Addressing these constraints will be essential to developing future traffic flow forecasting systems.

5. Industry Applications

Graph Neural Network (GNN) application in traffic flow forecasting extends beyond the purview of research; today, it is also making its presence felt in actual industrial applications. Scalable and accurate traffic forecasting instruments hold a very important place in urban mobility, logistics planning, navigation services, and the creation of autonomous vehicles.

It might be the most theatrical show of all in the actual world: Google Maps, which uses deep learning methods, including graph models, to forecast traffic flow and best routes. Google Maps leverages current location information from Android users and pools that information to create a graph of traffic wherein roads and intersections are nodes and lines. Although all the technicalities have not been disclosed by Google, it has been known that they utilize spatial-temporal relations like those utilized by structures utilized in making traffic predictions with the assistance of GNN (Google AI Blog, 2020).

Another colossus company, Uber, leverages traffic forecasting algorithms to enhance estimation of Estimated Time of Arrival (ETA) ride-sharing predictions. Uber's scenario is required to face highly dynamic patterns of traffic within urban settings and, therefore, is best suited for graph-based modelling. Uber publicly published a deep learning-based system called DeepETA in 2018 that employs road network models and patterns at a significant scale. Experiments suggest more up-to-date versions of such systems have since been built from the idea of graph neural networks with a focus on enhancing representations for road connectivity as well as time-dependent traffic flows (Ke et al., 2021).

Graph traffic forecasting systems also attract major investments under smart city initiatives. For example, the metropolis of Singapore has employed intelligent transport systems that integrate data from thousands of sensors, cameras, and GPS location-tracking devices to track and forecast traffic flow in real time. Singapore's Land Transport Authority (LTA) utilizes predictive analytics to schedule traffic lights, recommend diversions, and dynamically manage traffic crashes. Although technical specifics are not revealed, the application of graph-structured data points towards following GNN methodologies (LTA Singapore, 2021).

In autonomous driving, companies like Tesla and Waymo are exploring graph models to decide and plan routes. Forecasting traffic pattern movement around the vehicle is crucial for safety. GNNs enable these systems to sense relationships among vehicles, pedestrians, and infrastructure around them by representing the environment as graphs that dynamically change.

In logistics and supply chain management, businesses like Amazon and DHL are also beginning to introduce traffic forecasting into route algorithms. They will reduce transportation costs and improve customer satisfaction by anticipating congestion and optimizing routes in real-time.

These applications demonstrate that traffic flow prediction with GNN is not just a theoretical endeavor but a technology leading innovation in many industries. As the application of smart mobility technology continues to grow, the application of graph neural networks for powering real-time traffic prediction systems will grow even further.

6. Challenges and Limitations

Although the history of Graph Neural Networks (GNNs) in forecasting traffic flow has progressed a long way, there are challenges and problems. These need to be solved so that GNN models can be applied to real-world transport networks.

6.1 Data Sparsity and Incompleteness

There should be rich historical data and contemporaneous data for reliable traffic predictions. The majority of cities, particularly cities in the developing world, do not have the infrastructure to gather valuable traffic data. Even if cities are thoroughly surveyed, delays in communication, sensor faults, and loss of data might lead to lost or incomplete data. GNNs whose performance heavily depends upon graph node relationships and attributes become vulnerable to incomplete or noisy data and thus depreciate the predictability (Cui et al., 2019).

6.2 Scalability to Large-Scale Networks

As the number of cities and sensors, and cars increases, the traffic network graph grows exponentially quickly. All existing GNN models, such as DCRNN and STGCN, encounter computational issues when dealing with very large graphs. Training time and inference both become significantly longer, and memory requirements become a bottleneck (Wu et al., 2020). Whereas methods like GraphSAGE attempt to overcome the same by employing neighborhood sampling, further room for improvement exists so that GNNs may potentially work at the city-level, or the scale of a country's traffic network.

6.3 Dynamic and Non-Stationary Traffic Patterns

Traffic flow dynamics are very sensitive and subject to sudden changes due to accidents, adverse weather, social gatherings, and road maintenance. Static network topologies assumed by most such models as STGCN, cannot handle these emergent changes optimally. Dynamic graph learning techniques such as Graph WaveNet and STDGAT do provide relief but come with this additional drawback, with added complexity and a requirement for highly advanced machinery for emergent changes in graph networks in real time (Song et al., 2020).

6.4 Interpretability and Transparency

Most GNN-based traffic models are black-box models with poor interpretability. In major applications like autonomous driving and traffic control, it is significant to explain why a model predicts a certain traffic condition. Although attention-based models like ASTGCN enhance transparency by pointing out important nodes and time steps, interpretable GNN models remain an open research topic (Guo et al., 2019).

6.5 Privacy and Ethical Considerations

Traffic prediction models usually operate with real-time location data harvested from personal mobile devices and vehicles. Ensuring model accuracy while ensuring data privacy is a main concern. Approaches such as federated learning and differential privacy are being explored to address the issues, but integrating them into GNN-based systems without sacrificing accuracy remains challenging (Wu et al., 2020).

In summary, despite the unprecedented potential shown by GNNs in forecasting traffic flow, the challenges of data quality restrictions, scalability, flexibility, explainability, and privacy will be crucial to enable larger real-world applications.

7. Future Directions

As much as Graph Neural Networks (GNNs) have been achieving strong advances in traffic prediction, constant innovation still pushes the limits on what the models are able to achieve. A variety of promising research directions that arise would solve existing constraints and pave the way for even more powerful, larger, and more thorough traffic prediction models.

7.1 Dynamic Graph Modeling

One of the key lines of the future would be developing more sophisticated dynamic graph modeling methodology. Traffic networks are dynamic networks with connection and flow behavior that dynamically change with time as a result of incidents, road construction, and seasonal usage. Future models could involve real-time graph evolution where the graph structure continually changes in real time from real-time sensor observations. Methods such as source code in GitHub repository Temporal Graph Networks (TGN) (Rossi et al., 2020) can be trained on dynamic graphs and utilized for traffic applications.

7.2 Federated Learning for Privacy Preservation

With growing privacy concerns among users, federated learning provides a viable means of training GNN-based traffic models without centralizing personal data. Traffic information in a federated setting would remain on local devices or edge servers, and updates to the model alone would be exchanged for global training. Combining federated learning with GNNs would enable firms and governments to deploy robust traffic forecasting systems without violating GDPR (Li et al., 2020).

7.3 Multimodal Data Integration

Traffic behavior is affected by numerous external factors unrelated to past vehicle paths, such as weather, social events, and even human behavior. Multimodal GNNs could be the subject of future research that would enable traffic sensor data to be combined with external data sources such as weather reports, event calendars, or even social media trends. This would enable models to learn richer contextual signals, enhancing short- and long-term traffic prediction (Zhang et al., 2020).

7.4 Lightweight and Efficient GNN Architectures

Scalability becomes a problem with ever-larger and larger cities and increasingly large volumes of data. The next research would be developing light-weight GNN models involving low computation and memory expense, but with no loss in accuracy. Graph sampling, pruning, and quantization can be made to be beneficial to utilize in traffic prediction models to be run in real-time on low-power devices like edge computing hardware in smart intersections (Wu et al., 2020).

7.5 Explainable and Trustworthy GNNs

Because traffic forecasting is growing increasingly important to urban public amenities and security, it will need to render the GNNs interpretable and reliable. Experimentations of Explainable AI (XAI) to use with GNNs can make it possible for the models to predict traffic movement while also explaining how they came up with the prediction, making urban planners, managers, and commuters more confident (Ying et al., 2019).

8. Conclusion

But serious problems need to be resolved. Problems of sparsity of data, scalability of models, dynamic traffic adaptation, interpretability, and protection of privacy need to be resolved to achieve the full potential of GNNs in real-world applications. The area is at a stage of ongoing development, and new directions of research are moving towards dynamic graph modeling, integration of federated learning, multimodal fusion of data, lightweight architecture, and explainable GNN systems.

In general, Graph Neural Networks have emerged as the new norm in traffic flow prediction as a trustworthy model, surpassing the conventional models with greater precision and adaptability. With technological advancements driving it further, GNNs are sure to become an integral component of the smart transport of the future. Its evolution not only anticipates improved traffic prediction but also directly supports worldwide overall goals of fewer city congestion, more environmental gains, and building the urban way of life in the world.

References

Ahmed, M. S., & Cook, A. R. (1979). Analysis of freeway traffic time-series data by using Box-Jenkins techniques. *Transportation Research Record*, (722).

Bronstein, M. M., Bruna, J., LeCun, Y., Szlam, A., & Vandergheynst, P. (2017). Geometric deep learning: Going beyond Euclidean data. *IEEE Signal Processing Magazine, 34*(4), 18–42. https://ieeexplore.ieee.org/document/7974879

Google AI Blog. (2020, October 15). *How Google Maps predicts traffic and routes.* https://blog.google/products/maps/google-maps-101-how-ai-helps-predict-traffic-and-determine-routes/

Guo, S., Lin, Y., Feng, N., Song, C., & Wan, H. (2019). Attention-based spatial-temporal graph convolutional networks for traffic flow forecasting. In *Proceedings of the AAAI Conference on Artificial Intelligence* (Vol. 33, pp. 922–929). https://ojs.aaai.org/index.php/AAAI/article/view/3881

Hamilton, W. L., Ying, R., & Leskovec, J. (2017). Inductive representation learning on large graphs. In *Advances in Neural Information Processing Systems* (pp. 1024–1034).

Hochreiter, S., & Schmidhuber, J. (1997). Long short-term memory. *Neural Computation, 9*(8), 1735–1780. https://direct.mit.edu/neco/article-abstract/9/8/1735/6109/Long-Short-Term-Memory?redirectedFrom=fulltext

Ke, J., Zheng, H., Yang, H., & Chen, X. (2021). Deep learning for transportation: A survey. *IEEE Transactions on Intelligent Transportation Systems, 22*(6), 3133–3148. https://www.researchgate.net/publication/334660151_Deep_Learning_for_Intelligent_Transportation_Systems_A_Survey_of_Emerging_Trends

Kipf, T. N., & Welling, M. (2017). Semi-supervised classification with graph convolutional networks. In the *International Conference on Learning Representations (ICLR)*. https://arxiv.org/abs/1609.02907

Land Transport Authority Singapore. (2021). *Smart mobility 2030: Intelligent transport systems strategic* https://www.lta.gov.sg/content/dam/ltagov/getting_around/driving_in_singapore/intelligent_transport_systems/pdf/smartmobility2030.pdf

Li, T., Sahu, A. K., Talwalkar, A., & Smith, V. (2020). Federated learning: Challenges, methods, and future directions. *IEEE Signal Processing Magazine, 37*(3), 50–60. https://ieeexplore.ieee.org/document/9084352

Li, Y., Yu, R., Shahabi, C., & Liu, Y. (2018). Diffusion convolutional recurrent neural network: Data-driven traffic forecasting. In the *International Conference on Learning Representations (ICLR)*. https://arxiv.org/abs/1707.01926

Song, C., Lin, Y., Guo, S., & Wan, H. (2020). Spatial-temporal dynamic graph attention network for traffic flow forecasting. In *Proceedings of the AAAI Conference on Artificial Intelligence* (Vol. 34, pp. 922–929). https://ojs.aaai.org/index.php/AAAI/article/view/5425

Veličković, P., Cucurull, G., Casanova, A., Romero, A., Liò, P., & Bengio, Y. (2018). Graph attention networks. In the *International Conference on Learning Representations (ICLR)*. https://arxiv.org/abs/1710.10903

GitHub Link: https://github.com/JahnaviPokala/jpokala_64061/tree/main/Final%20Project